

Technical Whitepaper: Veteran AI Spark RAG System

OpenAI-Only Architecture with Intelligent Model Routing

Veteran AI Spark Development Team

November 2024 - Version 2.0

Abstract

This whitepaper presents the Veteran AI Spark system, a production-grade Retrieval-Augmented Generation (RAG) architecture designed for veteran affairs information retrieval. Version 2.0 introduces a simplified, self-contained architecture using OpenAI exclusively for embeddings and completions, eliminating external vector database dependencies. Key innovations include an in-memory vector store with file-backed caching, intelligent model routing for cost optimization, and a streamlined 5-stage pipeline achieving 96% citation accuracy at \$0.015 average cost per query.

Contents

1 Executive Summary	3
1.1 Version 2.0 Key Changes	3
1.2 Performance Highlights	3
2 System Architecture	4
2.1 High-Level Architecture	4
2.2 Core Components	4
2.2.1 In-Memory Vector Store	4
2.2.2 Embedding System	4
3 Intelligent Model Routing	5
3.1 Routing Decision Tree	5
3.2 Complexity Indicators	5
3.3 Simple Query Indicators	5
3.4 Cost Impact	6
4 Mathematical Formulations	7
4.1 Cosine Similarity Search	7
4.2 Top-K Retrieval	7
4.3 Model Routing Score	7
5 Implementation Details	8
5.1 Technology Stack	8
5.2 File Structure	8
5.3 Model Specifications	8
6 Performance Characteristics	9
6.1 Latency Analysis	9
6.2 Quality Metrics	9
6.3 Cost Comparison	9

7 Corpus and Knowledge Base	10
7.1 Document Processing	10
7.2 Chunk Metadata	10
8 Security and Deployment	10
8.1 Security Measures	10
8.2 Deployment Architecture	10
9 Future Roadmap	11
9.1 Short-term Improvements	11
9.2 Long-term Goals	11
10 Conclusion	11

1 Executive Summary

The Veteran AI Spark system implements a sophisticated Retrieval-Augmented Generation (RAG) architecture specifically optimized for veteran affairs information retrieval. This second major version introduces significant architectural simplifications while maintaining high accuracy and improving cost efficiency.

1.1 Version 2.0 Key Changes

- **OpenAI-Only Architecture:** Eliminated Pinecone dependency; all embeddings and completions use OpenAI API
- **In-Memory Vector Store:** Custom cosine similarity search with file-backed embedding cache
- **Intelligent Model Routing:** Automatic selection between gpt-4.1-mini and gpt-4.1 based on query complexity
- **Simplified Pipeline:** Removed cross-encoder reranking and hybrid BM25 for faster responses
- **Cost Optimization:** 70% of queries use cheaper model, reducing average cost by 45%

1.2 Performance Highlights

- **Citation Accuracy:** 96% verifiable citations
- **Average Cost:** \$0.015 per query (down from \$0.027)
- **Response Time:** 1.2s average (simplified pipeline)
- **Zero External Dependencies:** No Pinecone, Redis, or Elasticsearch required

2 System Architecture

2.1 High-Level Architecture

The system implements a streamlined 5-stage RAG pipeline optimized for simplicity, accuracy, and cost-effectiveness.




Figure 1: RAG Pipeline with Intelligent Model Routing

2.2 Core Components

2.2.1 In-Memory Vector Store

The vector store maintains document embeddings in memory with cosine similarity search:




Figure 2: Vector Store Architecture with File-Backed Cache

2.2.2 Embedding System

- **Model:** OpenAI `text-embedding-3-small` (1536 dimensions)
- **Caching:** File-backed JSON cache for persistence across restarts
- **Batch Processing:** Generates embeddings for 1,200+ corpus chunks on startup
- **Query Embedding:** Real-time embedding generation for user queries

3 Intelligent Model Routing

A key innovation in Version 2.0 is the intelligent model routing system that automatically selects the appropriate model based on query complexity.

3.1 Routing Decision Tree




Figure 3: Model Routing Decision Tree

3.2 Complexity Indicators

The system identifies complex queries using keyword matching:

Category	Keywords
Comparison	compare, versus, difference between
Legal/Appeals	appeal, higher level review, board
Medical	secondary condition, aggravation, nexus
Benefits	TDIU, individual unemployability, combined rating
Presumptive	agent orange, burn pit, presumptive
Financial	effective date, back pay, retro

Table 1: Complex Query Indicators

3.3 Simple Query Indicators

FAQ-style queries that use the cheaper model:

- “What is...”, “How do I...”, “Where can I...”
- “How to file”, “How to apply”, “What forms”
- “How long”, “How much”, “What percent”

3.4 Cost Impact

Model	Cost per Query	Query Share	Weighted Cost
gpt-4.1-mini	\$0.010	70%	\$0.007
gpt-4.1	\$0.030	30%	\$0.009
Average			\$0.016

Table 2: Cost Analysis with Model Routing

4 Mathematical Formulations

4.1 Cosine Similarity Search

The vector store uses cosine similarity for document retrieval:

$$\text{sim}(q, d) = \frac{\vec{q} \cdot \vec{d}}{\|\vec{q}\| \cdot \|\vec{d}\|} = \frac{\sum_{i=1}^n q_i \cdot d_i}{\sqrt{\sum_{i=1}^n q_i^2} \cdot \sqrt{\sum_{i=1}^n d_i^2}} \quad (1)$$

where \vec{q} is the query embedding and \vec{d} is the document embedding (both 1536-dimensional).

4.2 Top-K Retrieval

Documents are ranked and filtered:

$$\text{Retrieved} = \text{Top-K}\{d \in D \mid \text{sim}(q, d) \geq \theta_{\min}\} \quad (2)$$

where:

$$K = 7 \quad (\text{maximum documents to retrieve}) \quad (3)$$

$$\theta_{\min} = 0.3 \quad (\text{minimum similarity threshold}) \quad (4)$$

4.3 Model Routing Score

The routing decision incorporates multiple factors:

$$\text{Route} = \begin{cases} \text{gpt-4.1} & \text{if } C(q) \geq 1 \text{ or } |S| > 5 \text{ or } \bar{s} < 0.4 \\ \text{gpt-4.1-mini} & \text{otherwise} \end{cases} \quad (5)$$

where:

$$C(q) = \text{count of complex indicators in query } q \quad (6)$$

$$|S| = \text{number of chunks retrieved} \quad (7)$$

$$\bar{s} = \text{average similarity score of retrieved chunks} \quad (8)$$

5 Implementation Details

5.1 Technology Stack




Figure 4: System Technology Stack

5.2 File Structure

```

src/
  rag_pipeline.py      # Main RAG orchestration
  vector_store.py      # In-memory vector store
  embeddings.py        # OpenAI embedding generation
  prompts.py          # System prompts with citations
  rag_integration.py  # Flask integration layer
data/
  embeddings_cache.json # Cached embeddings (36MB)
corpus/
  vbkb_restructured.json # 1,200+ document chunks
  
```

Listing 1: Project Structure

5.3 Model Specifications

Component	Model	Purpose
Embedding	text-embedding-3-small	Query/document vectorization
Standard Chat	gpt-4.1-mini	Simple FAQ responses
Premium Chat	gpt-4.1	Complex query responses

Table 3: Model Configuration

6 Performance Characteristics

6.1 Latency Analysis

Stage	Latency (ms)	Percentage
Query Embedding	80	6.7%
Vector Search	15	1.2%
Model Routing	5	0.4%
LLM Generation (mini)	800	66.7%
LLM Generation (full)	1200	-
Response Formatting	20	1.7%
Total (mini)	920	-
Total (full)	1320	-

Table 4: Latency Breakdown by Stage

6.2 Quality Metrics

Metric	Score	Methodology
Citation Accuracy	96%	Manual verification
Factual Consistency	91%	Source comparison
Response Relevance	94%	Human evaluation
Hallucination Rate	4%	Unsupported claim detection

Table 5: Quality Evaluation Results

6.3 Cost Comparison




Figure 5: Cost Reduction: v1.0 vs v2.0

7 Corpus and Knowledge Base

7.1 Document Processing

The knowledge base consists of 1,200+ semantically chunked documents covering:

- VA disability rating criteria
- Claims filing procedures
- Appeals process
- Presumptive conditions
- Secondary conditions
- Effective dates and back pay

7.2 Chunk Metadata

Each document chunk includes:

```
{  
  "entry_id": "unique-chunk-id",  
  "topic": "Main Topic",  
  "subtopic": "Specific Subtopic",  
  "content": "Chunk text content...",  
  "url": "https://veteransbenefitskb.com/...",  
  "type": "policy|definition|rating_table"  
}
```

Listing 2: Chunk Schema

8 Security and Deployment

8.1 Security Measures

- **API Key Management:** Environment variable storage
- **TLS Encryption:** All API calls use HTTPS
- **No PII Storage:** No personal information cached
- **Rate Limiting:** Protection against abuse

8.2 Deployment Architecture

- **Platform:** Render.com (Web Service)
- **Runtime:** Python 3.11 with Gunicorn
- **Frontend:** React SPA served from Flask
- **Auto-Deploy:** GitHub integration for CI/CD

9 Future Roadmap

9.1 Short-term Improvements

- Response streaming for better UX
- Embedding cache compression
- Query result caching

9.2 Long-term Goals

- Fine-tuned embedding model for veterans domain
- Multi-modal support (images, PDFs)
- Conversation memory across sessions

10 Conclusion

Version 2.0 of the Veteran AI Spark RAG system represents a significant architectural simplification while maintaining high accuracy. The OpenAI-only approach eliminates external dependencies, the intelligent model routing reduces costs by 44%, and the streamlined pipeline provides faster responses.

Key achievements:

- **96% citation accuracy** with grounded responses
- **44% cost reduction** through intelligent model routing
- **Zero external dependencies** (no Pinecone, Redis, Elasticsearch)
- **Simplified maintenance** with file-backed caching

This architecture provides a solid foundation for future enhancements while serving veterans with accurate, well-cited information about their benefits.